

Explainable AI

Denis Nedić

dn98183@student.uni-lj.si

1. Model and LIME Setup

To understand the behavior of the black-box model, a Gradient Boosting Regressor was trained on the provided dataset. The categorical feature x_3 was properly handled using one-hot encoding prior to training to avoid introducing false ordinality assumptions. The Local Interpretable Model-Agnostic Explanations (LIME) framework was then implemented from scratch. For the surrogate model, Lasso regression was employed to inherently perform feature selection, yielding sparser and more interpretable local explanations.

2. Local Explanations & Feature Importance

The custom LIME algorithm was evaluated on the four specified instances. Across all instances, the explanation assigned a significant positive weight to x_1 (ranging from 2.02 to 3.14), identifying it as the overwhelmingly dominant predictive feature in these local neighborhoods.

Thanks to the Lasso regularization, the weight for x_2 was reduced to exactly 0.00 in all cases, and x_3 yielded either exactly 0.00 or a negligible contribution (~ 0.02). This sparse explanation is faithful to the true behavior of the black-box model. For example, comparing Instance 1 ($x_3 = 1$) and Instance 3 ($x_3 = 2$), the black-box predictions are identical (1.7449), confirming that x_3 has essentially no local effect. Similarly, modifying x_2 between Instance 3 and Instance 4 changed the black-box prediction by a mere 0.0085, verifying LIME's conclusion that x_2 is locally irrelevant.

Lastly, comparing the Black-Box and LIME predictions

demonstrates good local fidelity (e.g., 1.74 vs 1.80 for Instance 1). However, the gap in Instance 2 (0.77 vs 1.11) indicates that the true decision surface is highly non-linear in that specific region.


3. Implementation Choices & Justification

To ensure the robustness of the LIME explanations, a comprehensive sensitivity analysis was conducted on the core hyperparameters.


Neighborhood Size (N): The stability of the explanation relies on drawing sufficient samples. Figure 1(a) demonstrates the variance of the x_1 coefficient across multiple runs for varying N . The variance stabilizes significantly around $N = 1000$. Hence, $N = 1000$ was selected as the optimal trade-off between computational efficiency and explanation reliability.

Lasso Regularization (α): The regularization parameter dictates the sparsity of the explanation. Figure 1(b) plots the coefficient paths. Increasing α shrinks the weights towards zero. An $\alpha = 0.01$ was chosen because it successfully eliminates noise (forcing irrelevant features to zero) while preserving the valid, dominant contribution of x_1 .


Kernel Width: The exponential kernel width defines the locality of the neighborhood. Figure 1(c) illustrates the local fidelity (measured via weighted R^2 score) against different kernel widths. A width of 0.75 yields a peak fidelity score, ensuring the surrogate model accurately mimics the black-box within the relevant local region without overfitting to immediate noise or underfitting by incorporating too much of the global surface.



(a) Explanation Stability (N)



(b) Lasso Reg. Path (α)



(c) Local Fidelity (Kernel)

Figure 1. Parameter justification analysis for the LIME implementation.